

Planning Algorithms: An Overview

Zamboni Riccardo

e-Novia, Italy, ricc.zamboni@gmail.com

Keywords: Planning, MCTS, UCT, OLOP, UGape

Abstract

Planning algorithms are gaining increasing interest, as they are a powerful framework to deal with decision processes, and consist of a fruitful intersection between games and more general policy based scenarios. We consider the problem of planning in general stochastic and discounted environments. More precisely, the decision making problem consists in an exploration phase followed by a recommendation. First, the agent explores freely the set of possible sequences of actions, using a budget of n actions. Then the agent makes a recommendation on the first action to play. The goal of the agent is to find the best way to explore its environment (first phase) so that, once the available resources have been used, he is able to make the best possible recommendation on the action to play in the environment.

1 Introduction

Starting from a simple RL framework to make things more clear, a planning is conceived as when an agent uses some generative model of the environment to acquire more knowledge about the best policy to be followed, and then act accordingly. Planning algorithms are based on many different settings, and they can provide open-loop policies (independent from the state in which the agent is), and closed-loop policies, in which the actions are based on the states they lead to. We will take into account various settings for the planning framework, starting from the most common formulation of the problem.

2 Multi-Armed Bandits

2.1 Overview

Imagine you are in a casino facing multiple slot machines and each is configured with an unknown probability of how likely you can get a reward at one play. The question is: What is the best strategy to achieve highest long-term rewards?

In this post, we will only discuss the setting of having an infinite number of trials. The restriction on a finite number of trials introduces a new type of exploration problem. For instance, if the number of trials is smaller than the number of slot machines, we cannot even try every machine to estimate the reward probability (!) and hence we have to behave smartly w.r.t. a limited set of knowledge and resources (i.e. time).

A naive approach can be that you continue to playing with one machine for many many rounds so as to eventually estimate the “true” reward probability according to the law of large numbers. However, this is quite wasteful and surely does not guarantee the best long-term reward.

Now let's give it a scientific definition.

A Bernoulli multi-armed bandit can be described as a tuple of $\langle \mathcal{A}, \mathcal{R} \rangle$, where:

- We have K machines with reward probabilities, $\theta_1, \dots, \theta_K$
- At each time step t , we take an action a on one slot machine and receive a reward r .
- \mathcal{A} is a set of actions, each referring to the interaction with one slot machine. The value of action a is the expected reward, $Q(a) = E[r|a] = \theta$. If action a_t at the time step t is on the i -th machine, then $Q(a_t) = \theta_i$
- R is a **reward function**. In the case of Bernoulli bandit, we observe a reward r in a stochastic fashion. At the time step t , $r_t = R(a_t)$ may return reward 1 with a probability $Q(a_t)$ or 0 otherwise.

It is a simplified version of Markov decision process, as there is no state S , and it well describes many decision processes, as it is in the planning framework. The goal is to maximize the cumulative reward $\sum_{t=1}^T r_t$. If we know the optimal action with the best reward, then the goal is same as to minimize the potential regret or loss by not picking the optimal action. The optimal reward probability θ^* of the optimal action a^* is:

$$\theta^* = Q(a^*) = \max_{a \in \mathcal{A}} Q(a) = \max_{1 \leq i \leq K} \theta_i$$

Our loss function is the total regret we might have by not selecting the optimal action up to the time step T :

$$\mathcal{L}_T = E\left[\sum_{t=1}^T (\theta^* - Q(a_t))\right]$$

2.2 Upper Confidence Bounds and UCB1

An useful concept which is often taken into account when dealing with planning and Bandits is the one of the confidence bounds, and in particular the upper one. Random exploration gives us an opportunity to try out options that we have not

known much about. However, due to the randomness, it is possible we end up exploring a bad action which we have confirmed in the past (bad luck!). To avoid such inefficient exploration, one approach is to be optimistic about options with high uncertainty and thus to prefer actions for which we haven't had a confident value estimation yet. Or in other words, we favor exploration of actions with a strong potential to have a optimal value.

The **Upper Confidence Bounds (UCB)** algorithm measures this potential by an upper confidence bound of the reward value, $U^t(a)$, so that the true value is below with bound $Q(a) \leq Q^t(a) + U^t(a)$ with high probability. The upper bound $U^t(a)$ is a function of $N_t(a)$; a larger number of trials $N_t(a)$ should give us a smaller bound $U^t(a)$.

In UCB algorithm, we always select the greediest action to maximize the upper confidence bound:

$$a_t^{UCB} = \operatorname{argmax}_{a \in A} Q^t(a) + U^t(a)$$

Now, the question is how to estimate the upper confidence bound. If we do not want to assign any prior knowledge on how the distribution looks like, we can get help from **Hoeffding's Inequality**, a theorem applicable to any bounded distribution.

Let X_1, \dots, X_t be i.i.d. (independent and identically distributed) random variables and they are all bounded by the interval $[0, 1]$. The sample mean is $\bar{X}_t = \frac{1}{t} \sum_{\tau=1}^t X_\tau$. Then for $u > 0$, we have:

$$\mathcal{P}[E[X] > \bar{X}_t + u] \leq \exp -2tu^2$$

Given one target action a , let us consider:

- $r_t(a)$ as the random variables,
- $Q(a)$ as the true mean,
- $Q^t(a)$ as the sample mean,
- u as the upper confidence bound, $u = U_t(a)$

Then we have,

$$\mathcal{P}[Q(a) > Q^t(a) + U_t(a)] \leq \exp -2tU_t(a)^2$$

We want to pick a bound so that with high chances the true mean is below the sample mean + the upper confidence bound. Thus $\exp -2tU_t(a)^2$ should be a small probability. Let's say we are ok with a tiny threshold p :

$$e^{-2tU_t(a)^2} = p \text{ Thus, } U_t(a) = \sqrt{\frac{-\log p}{2N_t(a)}}$$

One heuristic is to reduce the threshold p in time, as we want to make more confident bound estimation with more rewards observed. Set $p = t^{-4}$ and we get **UCB1 Algorithm**:

$$U_t(a) = \sqrt{\frac{2 \log t}{N_t(a)}} \text{ and } a_t^{UCB1} = \operatorname{argmax}_{a \in A} Q(a) + \sqrt{\frac{2 \log t}{N_t(a)}}$$

3 Monte-Carlo Tree Search and UCT

In the development of planning algorithms, the approach based on the Monte-Carlo exploration has utterly gained attention. In general, Monte Carlo tree search (MCTS) algorithms employ strong (heuristic) search methods to identify the best available action in a given situation. They gain knowledge by simulating the given problem and thus require at least its generative model (also known as a forward, simulation, or sample model), which may be simpler than the complete transition model of a given task. The MCTS planning process incrementally builds an asymmetric search tree that is guided in the most promising direction by an exploratory action-selection policy, which is computed iteratively. It improves when increasing the number of iterations. An MCTS iteration usually consists of four consecutive phases: (1) selection of actions already memorized in the tree (descent from the root to a leaf node), (2) expansion of the tree with new nodes, (3) playout, i.e., selection of actions until a terminal state is reached, and (4) backpropagation of the feedback up the tree. A tree policy guides the selection phase and a default policy guides the playout phase.

Subsequently to the planning, the starting move leading most frequently to a win is played. The efficiency of Monte-Carlo search heavily depends on the way moves are sampled during the episodes. A pseudo-code for Monte-Carlo Search is showed in 1.

```

1: procedure MonteCarloSearch(position)
2:   repeat
3:     search(position, rootply)
4:   until Timeout
5:   return bestMove(position);

6: function search(position, depth)
7:   if GameOver then
8:     return GameResult
9:   end if
10:  m := selectMove(position, depth);
11:  v := -search(position after move m, depth + 1);
12:  Add entry (position, m, depth, v, ...) to the TT
13:  return v;

```

Figure 1. MCTS Pseudo-Code

UCT is a Monte-Carlo search algorithm with a specific randomized move selection mechanism based on the UCB1 algorithm, which can be exploited to find a meaningful way to select which move/action to perform, modelling the move selection problem as a separate multi-armed bandit for every (explored) internal node. In order to find the best move in the root, one has to determine the best moves in the internal nodes as well (at least along the candidate principal variations). Since the estimates of the values of the alternative moves rely on the estimates of the values of the (best) successor nodes, we must have small estimation errors for the latter ones. Hence the problem reduces to getting the estimation error decay quickly. In order to achieve this, the algorithm must balance between testing an alternative that looks currently the best (to obtain a precise estimate) and the exploration of other alternatives (to ensure that some good alternative is not missed), as guaranteed from the UCB1 framework. The pseudo-algorithm of the UCT is showed in 2.

```

1: function selectMove(position, depth)
2:  nMoves := # available moves in position
3:  nsum := 0 {nsum will contain # times the descendants of position are considered}
4:  for i := 1 to nMoves do
5:    Let tte[i] be the TT entry matching the ith descendant of position
6:    if the entry tte[i] is invalid then
7:      return random move in position
8:    end if
9:    nsum := nsum + tte[i].n {tte[i].n = # times the ith descendant is considered}
10: end for
11: maze := -∞;
12: for i := 1 to nMoves do
13:   if tte[i].n = 0 then
14:     v := +∞ {Give high preference to an unvisited descendant}
15:   else
16:     v := tte[i].value + √(2 * ln(nsum)/tte[i].n)
17:   end if
18:   if v > maze then
19:     maze := v
20:   end if
21:   Let m be the ith move
22: end for
23: return m

```

Figure 2. UCT Pseudo-Code

4 Temporal Differences Extensions and SARSA-UCT(λ)

It is possible to include temporal differences formalisms in the MCTS framework. In particular, it is possible to make assumptions on the missing values of playout states (and actions) in a similar way as assumptions can be made on the initial values $V_{init}(s)$ and $Q_{init}(s; a)$. We introduce the notion of a playout value function, which replaces the missing value estimates in the non-memorized part of the state space with playout values $V_{playout}(s)$ and $Q_{playout}(s; a)$. The focus on the non-memorizing part as well comes from the fact that this is a peculiarity of the MCTS framework, not represented in the RL one. Then, a TD back-up can be introduced in the usual MCTS case as shown in the pseudo-code 3.

```

procedure BACKUPTDERRORS(episode)
  δsum ← 0                                ▷ cumulative decayed TD error
  Vnext ← 0
  for i = LENGTH(episode) down to 1
    (s, R) ← episode(i)
    if s is in tree
      Vcurrent ← tree(s).V
    else
      Vcurrent ← Vplayout(s)                ▷ assumed playout value
    δ ← R + γVnext - Vcurrent                ▷ single TD error
    δsum ← λγδsum + δ                      ▷ decay and accumulate
    if s is in tree
      tree(s).n ← tree(s).n + 1             ▷ update value
      α ← 1 / (tree(s).n)                  ▷ example of MC-like step-size
      tree(s).V ← tree(s).V + αδsum
    Vnext ← Vcurrent                      ▷ remember the value from before the update

```

Figure 3. TD Backup Pseudo-Code

5 The other side of the coin: Best-Arm Identification and UGape

As we have seen planning methods make often use of a Monte-Carlo formalism. Yet MC backups might lead to an update with high variance, due to the long run of the simulation, moreover such methods often assume stationarity of the reward distributions over the tree, which might be not always the case. On the other side of the coin, in order to take the best decision at the root, it might be sufficient to be able to select simply the best action. This problem falls under the class of best-arm identification. A proper planning is actually going to balance between the two needs, of a long run outcome (Monte Carlo), and of the short-run outcome (Best-Arm selection). We will briefly

discuss the latter case. Unlike the standard multi-armed bandit problem, where the goal is to maximize the cumulative sum of rewards obtained by the planner, in this problem the planner is evaluated on the quality of the arm(s) returned at the end of the exploration phase. This problem has been studied in the literature from two different perspectives: fixed budget and fixed confidence, and the best choice can be given with some fixed confidence ϵ over a subset of the possible arms m , leading to the so-called (ϵ, m) -optimality. The definitions are the same as the standard multi-armed setting, with the exception of the extension to m -best-arm mean and simple regret of the actions in a set S with respect to the m -best arms:

$$\mu_{(m)} = \max_{k \in A} \mu_k, \quad r_m = \max_{k \in A} r_k = \mu_{(m)} - \min_{k \in S} \mu_k$$

In the best arm case, by assuming boundedness of the distributions of the rewards of each action, the **UGape** algorithm was shown to be particularly effective, both in a fixed budget (sampling efficiency) and fixed confidence (simple regret efficiency) case. One of the upsides of this algorithm is that its core structure is actually the same in both cases, and the selection function is the one adapted to the objective. After having played each arm at least once, the next choices follow a **selection function** based on the upper confidence concept. At each time step t , UGape first uses the observations up to time $t - 1$ and computes an index $B_k(t) = \max_{i \neq k}^m U_i(t) - L_k(t)$ for each arm $k \in A$, where

$$U_k(t) = \hat{\mu}_k(t-1) + \beta_k(t-1), \quad L_k(t) = \hat{\mu}_k(t-1) - \beta_k(t-1)$$

where $\beta_k(t-1)$ is a confidence interval specific for the fixed budget and fixed confidence cases, but it is pretty similar to the UCB case, and $U_k(t)$ and $L_k(t)$ are high probability upper and lower bounds on the mean of arm k , μ_k , after $t - 1$ rounds.

The index $B_k(t)$ is an upper-bound on the simple regret r_k of the k th arm. An additional index for a set S can be defined as $B_S(t) = \max_{i \in S} B_i(t)$. Similar to the arm index, B_S is also defined in order to upper-bound the simple regret r_S with high probability.

After computing the arm indices, UGape finds a set of m arms $J(t)$ with minimum upper-bound on their simple regrets. From $J(t)$, it computes two arm indices

$$u_t = \arg \max_{j \notin J(t)} U_j(t), \quad l_t = \arg \min_{i \in J(t)} L_i(t)$$

where in both cases the tie is broken in favor of the arm with the largest uncertainty β .

Arms l_t and u_t are the worst possible arm among those in $J(t)$ and the best possible arm left outside $J(t)$, respectively, and together they represent how bad the choice of $J(t)$ could be.

Finally, the algorithm selects and pulls the arm $I(t)$ as the arm with the large uncertainty β among the two, observes a sample $X_I(t)$ from the related distribution of reward of the arm, updates the mean and the number of the pulls for the selected arm, and so on.

This is repeated until a termination condition is met.

6 Open Loop Planning: OLOP

While the previous settings can be called closed loop settings, because the action depends on the states, there is another possible called open-loop planning, where the class of considered policies (i.e. sequences of actions) are only function of time (and not of the underlying resulting states). This open-loop planning is in general sub-optimal compared to the optimal (closed-loop) policy (mapping from states to actions). However, here, while the planning is open-loop (i.e. we do not take into consideration the subsequent states in the planning), the resulting general policy is closed-loop (since the chosen action depends on the current state). In such cases, “optimism in the face of uncertainty” has been shown as extremely effective (the most promising sequences of actions are explored first). In the case of open-loop planning, upper confidence bounds (UCBs) are assigned to all sequences of actions, and the exploration expands further the sequences with highest UCB. In this framework, it is critical to understand that the value of a sequence of action being defined as the sum of discounted rewards along the path, thus the rewards obtained along any sequence provides information, not only about that specific sequence, but also about any other sequence sharing the same initial actions. OLOP is designed to use this property as efficiently as possible, to derive tight upperbounds on the value of each sequence of actions. In particular, since the value of a sequence is the sum of discounted rewards, one would like to explore more intensively the sequences starting with actions that already yielded high rewards. OLOP proceeds as follows. It assigns upper confidence bounds (UCBs), called B-values, to all sequences of actions, and selects at each round a sequence with highest B-value. At time $m = 0$, the B-values are initialized to +1 to guarantee optimistic exploration. Then, after episode $m \geq 1$, the B-values are defined as follows: For any $1 \leq h \leq L$, for any $a \in A^h$, let

$$T_a(m) = \sum_{s=1}^m \mathcal{I}a_{1:h}^s = a$$

be the number of times we played a sequence of actions beginning with a . Now we define the empirical average of the rewards for the sequence a as:

$$\hat{\mu}_a(m) = \frac{1}{T_a(m)} \sum_{s=1}^m Y_h^s \mathcal{I}a_{1:h}^s = a$$

if $T_a(m) > 0$, and 0 otherwise. The corresponding upper confidence bound on the value of the sequence of actions a is by definition:

$$U_a(m) = \sum_{t=1}^h (\gamma^t \hat{\mu}_{a_{1:t}}(m) + \gamma^t \sqrt{\frac{2 \log M}{T_{a_{1:t}}(m)}}) + \frac{\gamma^{h+1}}{1-\gamma}$$

if $T_a(m) > 0$ and +1 otherwise. Now that we have upper confidence bounds on the value of many sequences of actions we can sharpen these bounds for the sequences by defining the B-values as:

$$B_a(m) = \inf_{1 \leq h \leq L} U_{a_{1:h}}(m)$$

At each episode with highest B-value, observes the rewards Y_t^m provided by the environment, and updates the B-values. At the end of the exploration phase, OLOP returns an action that has been the most played.